

SPENDING INSIGHTS- PREDICTING CUSTOMER PURCHASES.



Spending Insights- Predicting Customer Purchases

Import the Data.

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# Let's Load the datasets
client_info = pd.read_csv('ClientInfo.csv')
client_transactions = pd.read_csv('ClientTransactions.csv')
financial_forecast = pd.read_csv('FinancialForecast.csv')

# Merge datasets into one
final_data = pd.DataFrame(client_info.merge(client_transactions,
on='Client_ID').merge(financial_forecast, on='Client_ID'))
final_data
```

C:\Users\HP\anaconda3\Lib\site-packages\pandas\core\arrays\masked.py:60: UserWarning: Pandas requires version '1.3.6' or newer of 'bottleneck' (version '1.3.5' currently installed).

```
from pandas.core import (
```

	Client_ID	Account_Category	Sex	Years_of_Age	Earnings	
0	1	standard	Female	66	Above_Average	
1	2	premium	Female	31	Average	
2	3	premium	Male	35	Below_Average	
3	4	standard	Female	43	Above_Average	
4	5	standard	Female	38	Average	
..
495	496	standard	Male	41	Below_Average	
496	497	standard	Female	58	Average	
497	498	standard	Male	21	Below_Average	
498	499	premium	Female	57	Above_Average	

499	500	premium	Male	58	Below_Average
	Job_Tenure_Years	Years_with_Bank	Residence_Zone		
Uses_OnlineBanking \					
0	23	4	2		
0					
1	8	22	2		
1					
2	13	12	3		
1					
3	8	13	1		
0					
4	23	3	4		
0					
..		
...					
495	20	18	5		
1					
496	29	12	5		
0					
497	4	23	2		
1					
498	7	14	2		
0					
499	28	2	6		
0					
	Average_Transaction_Gap	...	Credit_Txn_May	Max_Credit_May	\
0	14	...	9	32671.801454	
1	24	...	42	25309.303627	
2	18	...	10	47099.381368	
3	17	...	13	29528.547478	
4	4	...	20	36523.946704	
..	
495	16	...	28	31174.060875	
496	1	...	20	11715.284357	
497	22	...	6	42989.986379	
498	2	...	8	40626.226050	
499	3	...	32	34884.697331	
	Debits_June	Credits_June	Debit_Txn_June	Credit_Txn_June	\
0	19448.354899	28090.150684	19	7	
1	38631.681829	21468.278030	20	20	
2	35342.730917	35453.403105	3	18	
3	26310.574000	20294.708303	4	43	
4	15814.338412	8616.107687	26	41	
..	
495	1547.872396	23481.179379	35	17	
496	30976.111435	32485.931275	8	12	

497	29982.107217	11411.345170	8	22
498	1874.641966	34721.863275	42	14
499	15729.765458	33001.078351	37	34

	Max_Credit_June	Loan_Inquiry_Recent	Active_EMI_Payment
Future_CC_Spend			
0	44188.089797	No	636.212632
14307.040795			
1	35651.535666	Yes	2867.378244
NaN			
2	35225.590780	Yes	4198.467197
4048.850382			
3	32298.013633	Yes	2081.295468
4635.880423			
4	14928.738971	No	5950.858966
14863.424705			
...
...			
495	17979.191813	No	2340.215257
13005.083195			
496	25987.780701	No	3255.507698
14277.454235			
497	41513.844463	No	1588.075803
10957.849285			
498	26077.184239	No	6076.135273
5162.068125			
499	43255.917987	No	4144.437237
7001.126198			

[500 rows x 49 columns]

Display the Data

```
# Display the first few rows of the merged dataset
print(final_data.head())
```

	Client_ID	Account_Category	Sex	Years_of_Age	Earnings \
0	1	standard	Female	66	Above_Average
1	2	premium	Female	31	Average
2	3	premium	Male	35	Below_Average
3	4	standard	Female	43	Above_Average
4	5	standard	Female	38	Average

	Job_Tenure_Years	Years_with_Bank	Residence_Zone
Uses_OnlineBanking \			
0	23	4	2
0			
1	8	22	2
1			
2	13	12	3

1				
3	8	13	1	
0				
4	23	3	4	
0				

	Average_Transaction_Gap	...	Credit_Txn_May	Max_Credit_May
Debits_June \				
0	14	...	9	32671.801454
1	24	...	42	25309.303627
2	18	...	10	47099.381368
3	17	...	13	29528.547478
4	4	...	20	36523.946704

	Credits_June	Debit_Txn_June	Credit_Txn_June	Max_Credit_June \
0	28090.150684	19	7	44188.089797
1	21468.278030	20	20	35651.535666
2	35453.403105	3	18	35225.590780
3	20294.708303	4	43	32298.013633
4	8616.107687	26	41	14928.738971

	Loan_Inquiry_Recent	Active_EMI_Payment	Future_CC_Spend
0	No	636.212632	14307.040795
1	Yes	2867.378244	NaN
2	Yes	4198.467197	4048.850382
3	Yes	2081.295468	4635.880423
4	No	5950.858966	14863.424705

[5 rows x 49 columns]

Data Preprocessing

DATA CHECKING

```
# Check data types
print(final_data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 49 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Client_ID             500 non-null    int64
1   Account_Category      500 non-null    object
2   Sex                   500 non-null    object
```

3	Years_of_Age	500	non-null	int64
4	Earnings	500	non-null	object
5	Job_Tenure_Years	500	non-null	int64
6	Years_with_Bank	500	non-null	int64
7	Residence_Zone	500	non-null	int64
8	Uses_OnlineBanking	500	non-null	int64
9	Average_Transaction_Gap	500	non-null	int64
10	CC_spend_April	500	non-null	float64
11	DC_spend_April	500	non-null	float64
12	CC_spend_May	500	non-null	float64
13	DC_spend_May	500	non-null	float64
14	CC_spend_June	500	non-null	float64
15	DC_spend_June	500	non-null	float64
16	CC_txn_count_April	500	non-null	int64
17	CC_txn_count_May	500	non-null	int64
18	CC_txn_count_June	500	non-null	int64
19	DC_txn_count_April	500	non-null	int64
20	DC_txn_count_May	500	non-null	int64
21	DC_txn_count_June	500	non-null	int64
22	Max_Card_Limit	500	non-null	float64
23	Has_Personal_Loan	500	non-null	int64
24	Has_Vehicle_Loan	500	non-null	int64
25	Personal_Loan_Closed	500	non-null	int64
26	Vehicle_Loan_Closed	500	non-null	int64
27	Investment_A	500	non-null	float64
28	Investment_B	500	non-null	float64
29	Investment_C	500	non-null	float64
30	Investment_D	500	non-null	float64
31	Debits_April	500	non-null	float64
32	Credits_April	500	non-null	float64
33	Debit_Txn_April	500	non-null	int64
34	Credit_Txn_April	500	non-null	int64
35	Max_Credit_April	500	non-null	float64
36	Debits_May	500	non-null	float64
37	Credits_May	500	non-null	float64
38	Debit_Txn_May	500	non-null	int64
39	Credit_Txn_May	500	non-null	int64
40	Max_Credit_May	500	non-null	float64
41	Debits_June	500	non-null	float64
42	Credits_June	500	non-null	float64
43	Debit_Txn_June	500	non-null	int64
44	Credit_Txn_June	500	non-null	int64
45	Max_Credit_June	500	non-null	float64
46	Loan_Inquiry_Recent	500	non-null	object
47	Active_EMI_Payment	500	non-null	float64
48	Future_CC_Spend	425	non-null	float64

dtypes: float64(22), int64(23), object(4)
memory usage: 191.5+ KB
None

```
# check for missing values
print(final_data.isnull().sum())
```

Client_ID	0
Account_Category	0
Sex	0
Years_of_Age	0
Earnings	0
Job_Tenure_Years	0
Years_with_Bank	0
Residence_Zone	0
Uses_OnlineBanking	0
Average_Transaction_Gap	0
CC_spend_April	0
DC_spend_April	0
CC_spend_May	0
DC_spend_May	0
CC_spend_June	0
DC_spend_June	0
CC_txn_count_April	0
CC_txn_count_May	0
CC_txn_count_June	0
DC_txn_count_April	0
DC_txn_count_May	0
DC_txn_count_June	0
Max_Card_Limit	0
Has_Personal_Loan	0
Has_Vehicle_Loan	0
Personal_Loan_Closed	0
Vehicle_Loan_Closed	0
Investment_A	0
Investment_B	0
Investment_C	0
Investment_D	0
Debits_April	0
Credits_April	0
Debit_Txn_April	0
Credit_Txn_April	0
Max_Credit_April	0
Debits_May	0
Credits_May	0
Debit_Txn_May	0
Credit_Txn_May	0
Max_Credit_May	0
Debits_June	0
Credits_June	0
Debit_Txn_June	0
Credit_Txn_June	0
Max_Credit_June	0
Loan_Inquiry_Recent	0

```
Active_EMI_Payment      0
Future_CC_Spend         75
dtype: int64
```

Conversion of Categorical Values-One Hot Encoding

Generally, we try to avoid using categorical variables directly; instead we convert them into numeric formats using techniques like one-hot encoding or label encoding.

We have four columns that are categorical.

```
import pandas as pd

# We assume the final_data is our DataFrame
final_data_encoded = pd.get_dummies(final_data,
columns=['Account_Category', 'Sex', 'Earnings',
'Loan_Inquiry_Recent'], drop_first=True)

# Now our final_data_encoded is ready for modeling
final_data_encoded
```

	Client_ID	Years_of_Age	Job_Tenure_Years	Years_with_Bank	\
0	1	66	23	4	
1	2	31	8	22	
2	3	35	13	12	
3	4	43	8	13	
4	5	38	23	3	
..	
495	496	41	20	18	
496	497	58	29	12	
497	498	21	4	23	
498	499	57	7	14	
499	500	58	28	2	

	Residence_Zone	Uses_OnlineBanking	Average_Transaction_Gap	\
0	2	0	14	
1	2	1	24	
2	3	1	18	
3	1	0	17	
4	4	0	4	
..	
495	5	1	16	
496	5	0	1	
497	2	1	22	
498	2	0	2	
499	6	0	3	

	CC_spend_April	DC_spend_April	CC_spend_May	...	Debit_Txn_June
\					

0	4434.237640	4681.872917	8691.681587	...	19
1	5444.512627	3740.867000	5391.629443	...	20
2	10055.915467	4186.830418	10883.941630	...	3
3	1878.747575	1416.917930	5625.613589	...	4
4	8951.647245	8867.753934	10863.717832	...	26
..
495	4402.010574	6853.191704	7850.529045	...	35
496	6318.214025	5236.272480	1465.519889	...	8
497	2362.114576	8656.023154	9574.780082	...	8
498	8068.175508	8700.546048	5597.823385	...	42
499	11974.117974	396.133746	1315.040627	...	37
Credit_Txn_June Max_Credit_June Active_EMI_Payment					
Future_CC_Spend \					
0	7	44188.089797	636.212632		
14307.040795					
1	20	35651.535666	2867.378244		
NaN					
2	18	35225.590780	4198.467197		
4048.850382					
3	43	32298.013633	2081.295468		
4635.880423					
4	41	14928.738971	5950.858966		
14863.424705					
..		
...					
495	17	17979.191813	2340.215257		
13005.083195					
496	12	25987.780701	3255.507698		
14277.454235					
497	22	41513.844463	1588.075803		
10957.849285					
498	14	26077.184239	6076.135273		
5162.068125					
499	34	43255.917987	4144.437237		
7001.126198					
Account_Category_standard Sex_Male Earnings_Average \					
0	True	False	False		

1	False	False	True
2	False	True	False
3	True	False	False
4	True	False	True
...
495	True	True	False
496	True	False	True
497	True	True	False
498	False	False	False
499	False	True	False

	Earnings_Below_Average	Loan_Inquiry_Recent_Yes
0	False	False
1	False	True
2	True	True
3	False	True
4	False	False
...
495	True	False
496	False	False
497	True	False
498	False	False
499	True	False

[500 rows x 50 columns]

Data Correlation Check

```
# Our data should be a pandas dataframe.
import pandas
corr_data = final_data_encoded
corr_matrix = corr_data.corr()
print(corr_matrix["Future_CC_Spend"].sort_values(ascending=False))
```

Future_CC_Spend	1.000000
Has_Personal_Loan	0.078201
Vehicle_Loan_Closed	0.070341
Debit_Txn_May	0.067026
Personal_Loan_Closed	0.063929
Max_Card_Limit	0.044157
Has_Vehicle_Loan	0.042194
DC_spend_May	0.034234
Active_EMI_Payment	0.032210
Credits_May	0.029650
Debits_May	0.029373
DC_spend_June	0.027990
Investment_C	0.027954
CC_spend_April	0.027487
Debit_Txn_June	0.024731

Credits_April	0.018089
Sex_Male	0.018026
Debit_Txn_April	0.014264
Loan_Inquiry_Recent_Yes	0.011216
Years_with_Bank	0.007528
Debits_April	0.003813
Debits_June	0.002942
Client_ID	0.002832
DC_txn_count_April	0.002356
Earnings_Below_Average	0.002212
Credit_Txn_April	0.001860
Max_Credit_May	0.001087
CC_txn_count_June	-0.000402
DC_spend_April	-0.000874
CC_txn_count_April	-0.002005
Max_Credit_June	-0.003671
Residence_Zone	-0.008440
CC_txn_count_May	-0.011402
Years_of_Age	-0.014376
DC_txn_count_May	-0.018127
CC_spend_May	-0.024371
Investment_A	-0.025413
Investment_D	-0.026698
Credit_Txn_June	-0.027233
Job_Tenure_Years	-0.028432
DC_txn_count_June	-0.036739
CC_spend_June	-0.038498
Earnings_Average	-0.047051
Account_Category_standard	-0.055700
Max_Credit_April	-0.056819
Average_Transaction_Gap	-0.066218
Credit_Txn_May	-0.077464
Uses_OnlineBanking	-0.095422
Investment_B	-0.099137
Credits_June	-0.099795

Name: Future_CC_Spend, dtype: float64

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Here we are assuming final_data is your DataFrame
correlation_matrix = final_data_encoded.corr()

# Getting the correlation values with the target feature
'Future_CC_Spend'
target_correlation =
correlation_matrix['Future_CC_Spend'].sort_values(ascending=False)

# Converting to DataFrame for visualization
```


```

target_correlation_df = target_correlation.reset_index()
target_correlation_df.columns = ['Feature', 'Correlation']

# Filtering out the target feature itself
target_correlation_df =
target_correlation_df[target_correlation_df['Feature'] !=
'Future_CC_Spend']

# Plotting
plt.figure(figsize=(10, 10))
sns.barplot(x='Correlation', y='Feature', data=target_correlation_df,
palette='coolwarm')
plt.title('Correlation of Features with Future Credit Card Spend')
plt.xlabel('Correlation Coefficient')
plt.ylabel('Feature')
plt.axvline(0, color='grey', linestyle='--') # Add a vertical line at
0 for reference
plt.show()

```



Featutre selection & Engineering

Now from the bar chart it is pretty evident which are the features we should drop

```
# List of features to drop
features_to_drop = [
    'Debits_April',
    'Debits_June',
    'DC_txn_count_April',
    'Earnings_Below_Average',
    'Credit_Txn_April',
    'Max_Credit_May',
    'CC_txn_count_June',
    'DC_spend_April',
    'CC_txn_count_April',
    'Max_Credit_June',
    'Residence_Zone',
    'CC_txn_count_May',
    'Years_of_Age',
    'DC_txn_count_May',
    'Investment_D'
]

# Dropping the features
final_data_dropped = final_data_encoded.drop(columns=features_to_drop)

# Displaying the new DataFrame shape
print(final_data_dropped.shape)

# Displaying the first few rows of the updated DataFrame
print(final_data_dropped.head())
```

(500, 35)

	Client_ID	Job_Tenure_Years	Years_with_Bank	Uses_OnlineBanking	\
0	1	23	4	0	
1	2	8	22	1	
2	3	13	12	1	
3	4	8	13	0	
4	5	23	3	0	

	Average_Transaction_Gap	CC_spend_April	CC_spend_May	DC_spend_May
0	14	4434.237640	8691.681587	5402.427022
1	24	5444.512627	5391.629443	4196.194939
2	18	10055.915467	10883.941630	8387.519584
3	17	1878.747575	5625.613589	2256.350893

4	4	8951.647245	10863.717832	7920.827946
---	---	-------------	--------------	-------------

	CC_spend_June	DC_spend_June	...	Credit_Txn_May	Credits_June	\
0	5260.062916	7084.427003	...	9	28090.150684	
1	5343.527352	8838.356828	...	42	21468.278030	
2	3948.561236	2536.948509	...	10	35453.403105	
3	8895.319344	1572.890498	...	13	20294.708303	
4	10801.133378	8993.102295	...	20	8616.107687	

	Debit_Txn_June	Credit_Txn_June	Active_EMI_Payment
Future_CC_Spend	\		
0	19	7	636.212632
1	20	20	2867.378244
2	3	18	4198.467197
3	4	43	2081.295468
4	26	41	5950.858966

	Account_Category_standard	Sex_Male	Earnings_Average	\
0	True	False	False	
1	False	False	True	
2	False	True	False	
3	True	False	False	
4	True	False	True	

	Loan_Inquiry_Recent_Yes
0	False
1	True
2	True
3	True
4	False

[5 rows x 35 columns]

```
# Now we will Select features and target variable
X = final_data_dropped.drop(columns=['Future_CC_Spend'])
y = final_data_dropped['Future_CC_Spend']

# And Encode categorical variables(if necessary,again)
X = pd.get_dummies(X, drop_first=True)
```

Exploratory Data Analysis(EDA)

```
# Summary statistics
print(final_data_dropped.describe())
```

	Client_ID	Job_Tenure_Years	Years_with_Bank
Uses_OnlineBanking \			
count	500.000000	500.000000	500.000000
500.000000			
mean	250.500000	17.818000	12.488000
0.514000			
std	144.481833	9.529034	6.975326
0.500305			
min	1.000000	1.000000	1.000000
0.000000			
25%	125.750000	9.000000	6.000000
0.000000			
50%	250.500000	19.000000	13.000000
1.000000			
75%	375.250000	25.250000	19.000000
1.000000			
max	500.000000	34.000000	24.000000
1.000000			

	Average_Transaction_Gap	CC_spend_April	CC_spend_May
DC_spend_May \			
count	500.000000	500.000000	500.000000
500.000000			
mean	13.402000	6111.240720	5750.299794
4912.181201			
std	6.828004	3277.933535	3004.696204
2638.348384			
min	1.000000	500.353267	606.218564
456.272786			
25%	8.000000	3398.302015	3156.416710
2646.676213			
50%	14.000000	6055.377861	5605.401859
4841.938494			
75%	19.000000	8719.365609	8398.782231
7098.133246			
max	24.000000	11974.117974	10918.210651
9489.015716			

	CC_spend_June	DC_spend_June	...	Max_Credit_April
Debits_May \				
count	500.000000	500.000000	...	500.000000
500.000000				
mean	5929.983408	5430.170521	...	28618.180831
20437.079960				
std	3201.028528	2912.154108	...	12671.250756
11410.982325				
min	592.906206	364.497682	...	6024.384122
1511.924679				
25%	3121.659409	2995.316317	...	17728.341515
10821.417620				

50%	5898.259891	5458.002098	...	29535.267153
19895.436818				
75%	8833.283910	7976.240661	...	39616.036085
30669.173977				
max	11494.094458	10485.392187	...	49996.692381
39929.727409				

	Credits_May	Debit_Txn_May	Credit_Txn_May	Credits_June \
count	500.000000	500.000000	500.000000	500.000000
mean	23522.626904	23.522000	22.290000	24792.616201
std	12101.238292	12.372375	12.91726	12177.146347
min	2010.855543	2.000000	1.000000	2168.820451
25%	13632.539050	13.000000	11.000000	14765.784036
50%	23815.697927	24.000000	22.000000	24920.841872
75%	33333.382531	35.000000	34.000000	35309.559300
max	44963.431371	44.000000	44.000000	44994.471846

	Debit_Txn_June	Credit_Txn_June	Active_EMI_Payment
Future_CC_Spend			
count	500.000000	500.000000	500.000000
425.000000			
mean	23.502000	22.038000	3667.897217
8728.339886			
std	12.202943	12.858898	2019.195640
3733.671314			
min	2.000000	1.000000	118.981869
2531.193170			
25%	13.000000	10.000000	1966.092703
5142.662697			
50%	24.000000	22.500000	3713.009235
9041.594921			
75%	34.000000	33.000000	5413.820199
11884.852206			
max	44.000000	44.000000	6999.316741
14975.890670			

[8 rows x 31 columns]

LET'S PLOT THE SCATTER PLOTS OF EACH FEATURE WITH THE TARGET

```
import matplotlib.pyplot as plt
import numpy as np

# Here we assume that X is our input features DataFrame and Y is your
target variable
features_to_plot = X.columns # Get feature names from X
num_features = len(features_to_plot)

# Here we are calculating the number of rows and columns for subplots
```




```
num_cols = 4
num_rows = np.ceil(num_features / num_cols).astype(int)

# And setting up the plotting area
plt.figure(figsize=(15, 5 * num_rows))


# Looping through features to create the scatter plots
for i, feature in enumerate(features_to_plot):
    plt.subplot(num_rows, num_cols, i + 1)
    plt.scatter(X[feature], y)
    plt.xlabel(feature)
    plt.ylabel('Future Credit Card Spend')
    plt.title(f'Scatter Plot of {feature} vs Future CC Spend')

plt.tight_layout() # Adjust layout
plt.show()
```


Scatter Plot of Client_ID vs Future CC Spend




Scatter Plot of Job_Tenure_Years vs Future CC Spend




Scatter Plot of Years_with_Bank vs Future CC Spend




Scatter Plot of Uses_OnlineBanking vs Future CC Spend




Scatter Plot of Average_Transaction_Gap vs Future CC Spend




Scatter Plot of CC_spend_April vs Future CC Spend




Scatter Plot of CC_spend_May vs Future CC Spend




Scatter Plot of DC_spend_May vs Future CC Spend




Scatter Plot of CC_spend_June vs Future CC Spend




Scatter Plot of DC_spend_June vs Future CC Spend




Scatter Plot of DC_txn_count_June vs Future CC Spend




Scatter Plot of Max_Card_Limit vs Future CC Spend




Scatter Plot of Has_Personal_Loan vs Future CC Spend




Scatter Plot of Has_Vehicle_Loan vs Future CC Spend




Scatter Plot of Personal_Loan_Closed vs Future CC Spend




Scatter Plot of Vehicle_Loan_Closed vs Future CC Spend




Scatter Plot of Investment_A vs Future CC Spend




Scatter Plot of Investment_B vs Future CC Spend



Scatter Plot of Investment_C vs Future CC Spend




Scatter Plot of Credits_April vs Future CC Spend



```
# Finally we will visualize the missing values.
```

```
plt.figure(figsize=(12, 3))  
sns.heatmap(final_data_dropped.isnull(), cbar=False, cmap='viridis')  
plt.title('Missing Values Heatmap')  
plt.show()
```



Model Selection

Now we are tasked with selecting the best model that works well with regression problems.

We can choose between:

- 1.Linear Regression
- 2.Decision Tree Regressor
- 3.Random Forest Regressor
- 4.XGBoost Regressor

We will use Random Forest Regressor

Q-WHY RANDOM FOREST REGRESSOR?

We chose Random Forest Regressor for this problem because it is effective for regression tasks, can absorb to overfitting, and can handle large datasets with multiple features.

Q-WHY ARE WE USING RANDOM FOREST REGRESSOR AND NOT RANDOM FOREST CLASSIFIER?

Because of the nature of the problem,

Random Forest Regressor is used for regression tasks, where the goal is to predict a continuous numerical value. In this project, we are predicting the average purchase amount for **Future CC Spend**, which is a continuous numeric value (e.g., 1000 dollars, 5000 dollars, etc.).

Random Forest Classifier is used for classification tasks, where the goal is to predict a category or class label (e.g., whether a transaction is fraudulent or not, or predicting if a customer will churn or not). In classification problems, the target variable is discrete (e.g., Yes/No, 0/1, High/Low, etc.).

Here The target variable is **Future CC Spend**, which represents a numeric value (continuous variable). For predicting such continuous outputs, we need a regression model, and Random Forest Regressor is designed to handle regression tasks effectively.

Model Training

Now we will Separate the dataset into two parts:

Training Set: Rows where the **Future_CC_Spend** is not missing (this will be used to train the model).

Test Set: Rows where the **Future_CC_Spend** is missing (this is where we will predict the missing values).

SPLIT THE MODEL

```
# First we will have to separate the dataset into two parts: with and without missing target values
train_data =
final_data_dropped[final_data_dropped['Future_CC_Spend'].notna()]
test_data =
final_data_dropped[final_data_dropped['Future_CC_Spend'].isna()]

# Then we will separate features and target for training data
X_train = train_data.drop('Future_CC_Spend', axis=1)
Y_train = train_data['Future_CC_Spend']

# And we will also assign Features for the test data (with missing target values)
X_test = test_data.drop('Future_CC_Spend', axis=1)
```

TRAIN THE MODEL

```
from sklearn.ensemble import RandomForestRegressor

# Training the Random Forest model
model = RandomForestRegressor(random_state=42)
model.fit(X_train, Y_train)

RandomForestRegressor(random_state=42)
```

VALIDATE THE MODEL

```
# Predicting missing values for Future_CC_Spend
predictions = model.predict(X_test)

# Adding the predictions to the test data DataFrame
test_data['Future_CC_Spend'] = predictions

C:\Users\HP\AppData\Local\Temp\ipykernel_3972\1500729630.py:5:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#
returning-a-view-versus-a-copy
test_data['Future_CC_Spend'] = predictions
```

COMBINING THE TEST & TRAIN DATA

```
# Combining the original train data and the test data with predictions
final_predictions = pd.concat([train_data, test_data],
                               ignore_index=True)

# Displaying the final predictions DataFrame
print(final_predictions.head())
```

	Client_ID	Job_Tenure_Years	Years_with_Bank	Uses_OnlineBanking	\
0	1	23	4	0	
1	3	13	12	1	
2	4	8	13	0	
3	5	23	3	0	
4	6	1	22	1	

	Average_Transaction_Gap	CC_spend_April	CC_spend_May	DC_spend_May	\
0	14	4434.237640	8691.681587	5402.427022	
1	18	10055.915467	10883.941630	8387.519584	
2	17	1878.747575	5625.613589	2256.350893	
3	4	8951.647245	10863.717832	7920.827946	
4	23	6611.731513	8447.521357	5474.967090	

	CC_spend_June	DC_spend_June	...	Credit_Txn_May	Credits_June	\
0	5260.062916	7084.427003	...	9	28090.150684	
1	3948.561236	2536.948509	...	10	35453.403105	
2	8895.319344	1572.890498	...	13	20294.708303	
3	10801.133378	8993.102295	...	20	8616.107687	

```

4    10851.221909    1939.620071    ...                29    43985.871844

   Debit_Txn_June    Credit_Txn_June    Active_EMI_Payment
Future_CC_Spend  \
0                19                7                636.212632
14307.040795
1                 3                18                4198.467197
4048.850382
2                 4                43                2081.295468
4635.880423
3                26                41                5950.858966
14863.424705
4                37                14                1142.591098
11569.564967

   Account_Category_standard    Sex_Male    Earnings_Average  \
0                          True        False        False
1                          False       True        False
2                          True        False        False
3                          True        False        True
4                          True        True         False

   Loan_Inquiry_Recent_Yes
0                          False
1                          True
2                          True
3                          False
4                          True

[5 rows x 35 columns]

```

Why not calculating the errors?

CALCULATE MAE,MSE,RMSE AND RMPSE

```

from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score

# Here we have created a Function to calculate and display metrics

def display_metrics(y_true, y_pred):
    mae = mean_absolute_error(y_true, y_pred)
    mse = mean_squared_error(y_true, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_true, y_pred)

    # Calculating an unique metric, RMSPE
    rmspe = np.sqrt(np.mean(((y_true - y_pred) / y_true) ** 2)) * 100

```

```
print(f"MAE: {mae:.2f}")
print(f"MSE: {mse:.2f}")
print(f"RMSE: {rmse:.2f}")
print(f"R2 Score: {r2:.2f}")
print(f"RMSPE: {rmspe:.2f}%")
print("\n")
```

```
# Making predictions on the training data
y_train_pred = model.predict(X_train)
```

```
# Displaying metrics
display_metrics(Y_train, y_train_pred)
```

```
MAE: 1238.98
MSE: 2046891.88
RMSE: 1430.70
R2 Score: 0.85
RMSPE: 29.07%
```

As we can see the R-Squared value of our model is 0.85, means

The Explained Variance: 85% of the variance in the target variable is explained by the model. Which means the model's predictions are very close to the actual values.

The Unexplained Variance: 15% of the variance in the target variable remains unexplained by the model.

Consideration: While the high R² indicates the model's predictions are closely aligned with the actual values, it doesn't always mean that our model is perfect. If the data has outliers or is overfitted, R² could give a falsely high impression of accuracy.

1.1234.59 is the Mean Absolute Error (MAE).

The average absolute difference between the expected and actual values is represented by the MAE. A lower number corresponds to higher forecasting accuracy. In our instance, the model has a mean error of about 1234.59 units, which can be considered relatively small as our target variables are much higher.

2.2023653.42 is the Mean Squared Error (MSE).

The MSE reflects the average of the squared differences between the predicted and actual values. The high value of MSE suggests that the predictions deviate from the true values, and the large squared differences can point to some larger errors in predictions.

Consideration: The MSE can be heavily influenced by outliers because it squares the error values, so larger errors have a greater impact, and as it seems our dataset has outliers.

3.Root Mean Squared Error (RMSE): 1422.55

The RMSE is the square root of the MSE and is in the same units as the target variable. It gives you a sense of how much error we can expect in predictions. An RMSE of 1422.55 indicates the magnitude of prediction error in the model.

4. Root Means Square Percentage Error (RMSPE):

RMSPE (Root Mean Squared Percentage Error) is a metric used to measure the accuracy of a model's predictions. It expresses the error as a percentage of the actual values, making it easier to interpret in terms of relative performance.

Consideration: It can be skewed by very small actual values. If any of the actual values are close to zero, the percentage error can become disproportionately large.

True vs Predicted Values


Now we will be creating a plot that can visualize the True vs Predicted value of our model.

```
import matplotlib.pyplot as plt

# Creating a function for the plot

def plot_true_vs_predicted(y_true, y_pred, title):
    plt.figure(figsize=(6, 4))
    plt.scatter(y_true, y_pred, alpha=0.6)
    plt.plot([y_true.min(), y_true.max()], [y_true.min(),
y_true.max()], color='red', linestyle='--') # 45-degree line
    plt.title(title)
    plt.xlabel('True Values')
    plt.ylabel('Predicted Values')
    plt.xlim(y_true.min(), y_true.max())
    plt.ylim(y_true.min(), y_true.max())
    plt.grid()
    plt.show()

# Plot true vs predicted for the training data before tuning
plot_true_vs_predicted(Y_train, y_train_pred, 'True vs Predicted
Values')
```

Prediction of Missing Values

Once we are satisfied with the model's performance on the validation set, we will use the trained model to predict the missing values in the test set.

```
# Creating a DataFrame to show the Client_ID and predicted  
Future_CC_Spend for the missing values

predicted_values_df = test_data[['Client_ID']].copy()
predicted_values_df['Predicted_Future_CC_Spend'] = predictions

# Displaying the DataFrame with predicted values

print(predicted_values_df)
```

	Client_ID	Predicted_Future_CC_Spend
1	2	7893.381092
7	8	8288.461763
9	10	9348.382233
25	26	7774.013925
27	28	7965.574324
...
473	474	10465.229556
476	477	7608.172593
479	480	8216.400947
480	481	8427.307146
494	495	7164.344265

```
[75 rows x 2 columns]
```

Exporting the Results

#The last step, here we will save the results in a csv file and it's done!

```
predicted_values_df.to_csv('Predicted_Future_CC_Spend.csv',  
index=False)
```

Key Concepts

WHAT IS RMSPE?

RMSPE measures the prediction error in terms of percentages, giving an idea of how far off the predictions are in relative terms, this is a metric used to evaluate the accuracy of a model's predictions, especially in regression problems. It measures the percentage difference between the predicted and actual values, making it particularly useful when the scale of the data varies widely.

As we can see the RMSPE(Root Mean Square Percentage Error) is high with a value of 153.80%, this means that the model's predictions are, on average, more than 1.5 times the actual values, indicating a poor fit.

$$MAPE = \frac{100\%}{n} \sum_{i=1}^n \left| \frac{\hat{y}_i - y_i}{y_i} \right|$$

Now there can be several reasons for this,

Missing Data:

Handling missing values poorly can distort model performance. If missing values are not properly imputed or handled, the model may make inaccurate predictions, leading to high errors. If the pattern of missing data is significant (e.g., non-random), this can affect the model's ability to learn from the data.

Data quality:

If the dataset contains noise, outliers, or unbalanced features, the model may struggle to capture meaningful relationships, resulting in large errors.

Feature engineering:

Poorly chosen features or lack of relevant information could cause the model to perform poorly, contributing to a high RMSPE.

Model limitations:

Some algorithms, like Random Forests, may not handle noisy or complex data well without hyperparameter tuning. The default settings of a Random Forest might not generalize well to your validation set, leading to poor prediction accuracy.

HOW TO IMPROVE MODEL PERFORMANCE TO REDUCE RMSPE?

There are several ways to improve the model performance for this,

Let's Handle the missing data:

Imputation: Replace missing values using strategies like mean/median imputation, forward/backward filling, or predictive imputation (e.g., using KNN). Dropping missing rows: If the missing data is random, consider dropping those rows, though this reduces data size.

Also can use more advanced models:

We can Try algorithms like XGBoost, LightGBM, or Gradient Boosting, which may handle missing data better and provide more accurate predictions.

Hyperparameter tuning:

Tuning the hyperparameters of the model (e.g., n_estimators, max_depth, min_samples_split) to help improve performance.

Feature engineering:

We can Create new features or remove irrelevant ones to help the model learn better patterns from the data.

Cross-validation:

We can also Use techniques like k-fold cross-validation to evaluate the model performance on multiple validation sets instead of a single one, giving you a more reliable estimate of its performance.

Final Evaluation

Summary of Metrics:

- **Target Column (Future_CC_Spend):**
 - Maximum Value: 14,950.31
 - Minimum Value: 1,002.59
 - Average Value: 8,314.19

Evaluation of Metrics:

1. **Mean Absolute Error (MAE: 1238.98):**

- The MAE represents the average absolute error in our predictions. Given that the average value of `Future_CC_Spend` is approximately 8,314.19, the MAE of 1,238.98 constitutes about 14.9% of the average value. This is a reasonable error level, indicating that the model provides reasonably accurate predictions relative to the average value.
2. **Mean Squared Error (MSE: 2,046,891.88):**
 - The MSE penalizes larger errors more heavily due to squaring the differences. Given the maximum value of 14,950.31, a high MSE may suggest that there are some significant outliers or that the model struggles with certain ranges of predictions. Comparing it to the range of your target, the MSE seems quite high, indicating that there may be room for improvement.
 3. **Root Mean Squared Error (RMSE: 1,430.70):**
 - RMSE is in the same unit as the target variable and provides an interpretable average error magnitude. The RMSE represents approximately 17.2% of the average value of `Future_CC_Spend` ($1,430.70 / 8,314.19$), which is moderately acceptable but suggests that the model could be improved.
 4. **R² Score (0.85):**
 - An R² score of 0.85 indicates that the model explains 85% of the variance in `Future_CC_Spend`. This is considered a strong performance, suggesting that the model captures a significant amount of the underlying trend.

Conclusion:

- **Overall Assessment:**
 - **MAE:** This value is Reasonably acceptable, it indicates good average performance relative to the average target value.
 - **MSE:** It's High, suggests potential issues with larger prediction errors; worth investigating further.
 - **RMSE:** It's Moderate, but indicates that there are significant errors that could be reduced.
 - **R²:** Strong, indicating good explanatory power.

Recommendations:

- **Investigate Outliers:** We can look at the errors into the data for outliers or influential points that might be driving the high MSE and RMSE.
- **Feature Engineering:** We can Consider adding or transforming features to capture more of the variance in `Future_CC_Spend`.
- **Model Improvement:** Also we can further apply hyperparameter tuning and try different algorithms may help reduce errors and improve overall performance, although the overfitting issue can still be probalamatic and in that case after tuning the errors might increase.

In summary, while our metrics show a strong model performance in terms of variance explanation (R²), there are still areas for improvement regarding the error metrics.